

# Profiling an Artificial Neural Network implementation

Francesco Cicala


*Assignment 2*

---

## Abstract

In this assignment we have selected a non-trivial C code and profiled it by means of a popular profiling software, Valgrind; in particular, we used the tools Memcheck, Callgrind and KCacheGrind. The code is the implementation of a simple neural network solving the XOR problem. It was provided by GenANN (<https://github.com/codeplea/genann>), a minimal library used for generating and training feedforward artificial neural networks in C.

---



## 1. Memcheck

Valgrind memcheck is a tool to detect memory errors in an execution. Memory errors could be:

- accessing memory after it has been freed;
- incorrect freeing of heap memory;
- memory leaks, i.e., memory was allocated and not freed before the program termination.

Firstly, we compiled the script and executed it with the following commands:

```
gcc -c example1.c -Wall -Wextra
gcc -c genann.c -Wall -Wextra
gcc example1.o genann.o -o example1.x -pg -lm -Wall -Wextra
./example1.x
```

And the output of the execution was:

```
GENANN example 1.
Train a small ANN to the XOR function using backpropagation.
Output for [0, 0] is 0.
Output for [0, 1] is 1.
Output for [1, 0] is 1.
Output for [1, 1] is 0.
```

After that, we profiled the memory usage of the execution:

```
valgrind --verbose --show-leak-kinds=all --leak-check=full
--track-origins=yes --log-file=valgrind.out ./example1.x
```

The output was:

```
==17118== Memcheck, a memory error detector
==17118== Copyright (C) 2002–2015, and GNU GPL'd, by Julian Seward et al.
==17118== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==17118== Command: ./example1.x
==17118== Parent PID: 15821
==17118==
[...]
```

```
==17118== embedded gdbserver: reading from
/tmp/vgdb-pipe-from-vgdb-to-17118-by-fra-on-???
```

```

==17118== embedded gdbserver: writing to
/tmp/vgdb-pipe-to-vgdb-from-17118-by-fra-on-???
==17118== embedded gdbserver: shared mem
/tmp/vgdb-pipe-shared-mem-vgdb-17118-by-fra-on-???
==17118==
==17118== TO CONTROL THIS PROCESS USING vgdb (which you probably
==17118== don't want to do, unless you know exactly what you're doing,
==17118== or are doing some strange experiment):
==17118== /usr/lib/valgrind/../../bin/vgdb --pid=17118 ...command...
==17118==
[...]
```

```

==17118==
==17118== HEAP SUMMARY:
==17118==   in use at exit: 0 bytes in 0 blocks
==17118== total heap usage: 3 allocs, 3 frees, 11,924 bytes allocated
==17118==
==17118== All heap blocks were freed -- no leaks are possible
==17118==
==17118== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
==17118== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

In the Heap Summary we see that no memory errors were generated, that is, all the heap blocks that were allocated were also freed before the program termination.

## 2. Callgrind and KCacheGrind

By calling callgrind with argument the executable we considered in the last section, a file callgrind.out is generated. We analyzed this file by means of KCacheGrind, a graphic user interface. Below, the time elapsed in the different functions is shown:

|   |       |      |       |         |      |                           |
|---|-------|------|-------|---------|------|---------------------------|
| ■ | 99.60 | 0.81 | 1     | ■       | main |                           |
| ■ | 97.32 | ■    | 56.55 | 40 000  | ■    | genann_train              |
| ■ | 40.78 | ■    | 24.90 | 40 004  | ■    | genann_run                |
| ! | 10.26 | !    | 10.26 | 120 012 | ■    | genann_act_sigmoid_ca...  |
| ! | 9.65  |      | 2.81  | 80 008  | ■    | genann_act_hidden_ind...  |
|   | 4.82  |      | 1.40  | 40 004  | ■    | genann_act_output_ind...  |
|   | 1.40  |      | 1.40  | 40 004  | ■    | __memcpy_avx_unaligned    |
|   | 1.39  |      | 0.00  | 1       | ■    | genann_init               |
|   | 1.38  |      | 0.13  | 1       | ■    | genann_init_sigmoid_lo... |
|   | 1.25  |      | 0.21  | 4 096   | ■    | genann_act_sigmoid        |
|   | 1.04  |      | 0.14  | 4 096   | ■    | exp                       |

As we see, the most of the time is passed in the function `genann_train`, which is the method by which the parameters of the neural network are learned until their convergence. `genann_train` calls `genann_run`, which is the method which processes the output for a given train input, so that a loss value is computed by comparing the output and the expected result.

