

# Multicore/Multinode hands-on exercises

Francesco Cicala

*Assignment 5*

---

## Abstract

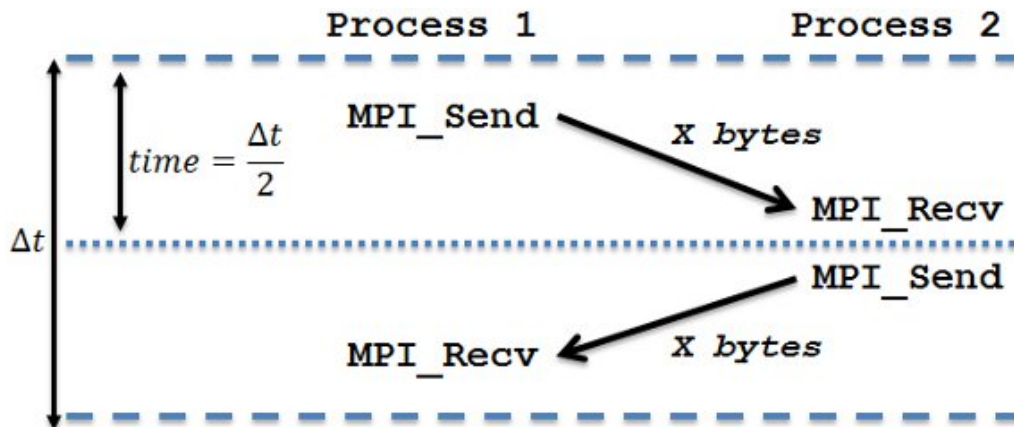
In this assignment we have tested internode, intercore performances on SISSA's Ulysses cluster. It has been done by means of some tools: the PingPong Intel's benchmark, the STREAM benchmark, hwloc, numactl and likwid. The full procedure is divided in three exercises: a measure of latency among MPI processes, bandwidth measures and the run of a program, nodeperf.c.

---

## 1. Latency and bandwidth measure with IMPI PingPong benchmark

The Intel MPI is a set of benchmarks for performance measurements regarding communication operations which are capable to characterize performances of a cluster system, network latency, throughput and the efficiency of an MPI implementation.

We have used the PingPong benchmark to assess the latency and the bandwidth during the exchange of packets with in a range of different sizes.



We show the results for two cases: intrasocket and intranode (intersocket).

### 1.1. Intrasocket

```
mpirun -np 2 hwloc--bind core:0 core:1 \  
/u/shared/programs/x86_64/intel/impi.5.0.1/bin64/IMB-MPI1 \  
PingPong -iter 100000 > run.dat
```

```
-----  
# Benchmarking PingPong  
# #processes = 2  
-----  
#bytes #repetitions      t[usec]    Mbytes/sec  
0      100000           0.17         0.00  
1      100000           0.19         5.11  
2      100000           0.18        10.52  
4      100000           0.18        21.04  
8      100000           0.18        42.08  
16     100000           0.18        84.11  
32     100000           0.20       148.98  
64     100000           0.26       239.02  
128    100000           0.26       468.39  
256    100000           0.28       880.25  
512     81920           0.34      1450.43  
1024    40960           0.40      2442.97  
2048    20480           0.55      3564.58  
4096    10240           0.83      4702.00  
8192     5120           1.42      5496.76  
16384    2560           2.61      5992.93  
32768    1280           4.53      6901.36  
65536     640           7.96      7846.97  
131072    320          14.19      8808.79  
262144    160          26.85      9312.14  
524288     80          49.54     10092.17  
1048576    40          94.79     10549.72  
2097152    20         180.85     11058.74  
4194304    10        353.96     11300.83
```

### 1.2. Intersockets

```
mpirun -np 2 hwloc--bind core:0 core:13 \  
/u/shared/programs/x86_64/intel/impi.5.0.1/bin64/IMB-MPI1 \  
PingPong -iter 100000 > run.dat
```

```

-----
# Benchmarking PingPong
# #processes = 2
-----

```

#bytes	#repetitions	t[usec]	Mbytes/sec
0	100000	0.47	0.00
1	100000	0.51	1.86
2	100000	0.51	3.73
4	100000	0.51	7.46
8	100000	0.51	14.92
16	100000	0.51	29.74
32	100000	0.52	58.86
64	100000	0.56	108.85
128	100000	0.57	213.40
256	100000	0.61	401.15
512	81920	0.80	613.53
1024	40960	1.00	979.03
2048	20480	1.30	1498.66
4096	10240	1.85	2113.61
8192	5120	3.24	2409.20
16384	2560	6.12	2552.17
32768	1280	11.46	2727.67
65536	640	11.59	5392.00
131072	320	21.08	5929.92
262144	160	39.62	6309.71
524288	80	78.19	6394.37
1048576	40	156.41	6393.27
2097152	20	315.23	6344.67
4194304	10	629.60	6353.20

### 1.3. Comparison

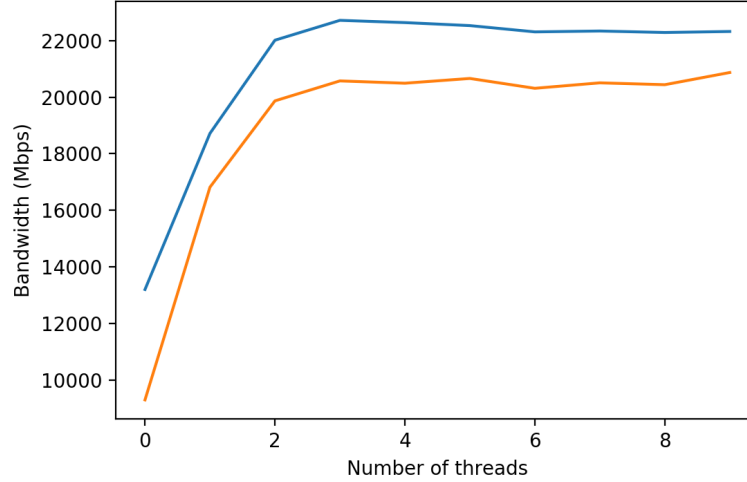
It is clear that optimizing code by being aware of the core on which each process is running can make a relevant difference in performances, when data exchange is involved. In fact, when the cores belong to the same socket, the latency is  $0.17\mu s$  and bandwidth reaches 11 Gbps by increasing the size of the packets. On the other hand, the intersockets case shows a significant increase in latency,  $0.47\mu s$ , and a bandwidth which is nearly a half of the previous one, approximately 6 Gbps.

## 2. Stream benchmark

Stream is a synthetic benchmark program that measures sustainable memory bandwidth. We have compiled it and executed the benchmark by means of the command:

```
(for i in `seq 1 10`; do OMP_NUM_THREADS=$i numactl --cpunodebind 0
--membind 0 ./stream.x; done) | grep "Triad:" | awk '{print $2}'
```

The second time, we have changed the value of `membind` to 1. In this way we have two measurements, each composed of ten different bandwidth measures (MB/s) for the number of threads in the range from 1 to 10.



We observe that, for the same number of threads, by using the local memory the performance is higher. This is easily explainable by considering that if the memory is local the data runs over shorter paths.

### 3. Nodeperf

We have firstly compiled the `nodeperf.c` file by using the Intel compiler, and then by means of the GNU compiler. The output is an executable, which computes matrix-matrix multiplication in order to measure the peak performance of an Ulysses node.

The command we used to invoke the Intel compiler is:

```
mpiicc -O2 -xHost -fopenmp -mkl nodeperf.c -o nodeperf.x
```

On the other hand, to compile it with the GNU compiler:

```
mpicc -fopenmp -o3 nodeperf.c -m64 -I${MKLROOT}/include -o
nodeperf.x -L${MKLROOT}/lib/intel64 -Wl,--no-as-needed
-lmkl_intel_lp64 -lmkl_sequential -lmkl_core -lpthread -lm -ldl
```

The executable produced by the Intel compiler has a performance of 442 GFlops, while the other one reached 27 GFlops.

The peak performance of a node is 48 GFlops. Therefore, the first executable reached the 98,6%, while the second one only the 6%. We conclude that the highly optimized compiler completely surpass the GNU one, and so, in general, the choice of the compiler can be hugely relevant in a performance perspective.